

Exam Guide

AWS Certified AI
Practitioner



350+ FREE Exam Questions



Yatharth Chauhan
Follow for more content

350+ Exam Preparation QnA (100% OFF)

The screenshot shows a Udemy course page for 'AWS Certified AI Practitioner AIF-C01: Exam Preparation'. The course is free, down from \$1,769. It has 4 days left at this price. The course includes 6 practice tests, assignments, access on mobile, and full lifetime access. It comes with a 30-day money-back guarantee. The course includes 350+ questions and is taught by Yatharth Chauhan in English. The course covers AWS AI services like Bedrock, SageMaker, Comprehend, RAG, embeddings, and governance.

IT & Software > IT Certifications > Generative AI (GenAI)

AWS Certified AI Practitioner AIF-C01: Exam Preparation

Master AWS AI services: Bedrock, SageMaker, Comprehend, RAG, embeddings, governance. 6 practice tests, 350+ Questions

Created by [Yatharth Chauhan](#)

English

What you'll learn

- ✓ Master AWS AI services (Bedrock, SageMaker, Comprehend) and pass the AIF-C01 certification exam
- ✓ Learn prompt engineering techniques including chaining, few-shot learning, and system prompts to improve AI model accuracy
- ✓ Understand foundation models, LLMs, and generative AI concepts; implement RAG systems for production applications
- ✓ Identify and mitigate AI bias, ensure responsible deployment, and understand governance requirements for production AI systems

Free ₹1,769 100% off
4 days left at this price!

Add to cart

Enroll now

30-Day Money-Back Guarantee

This course includes:

- 6 practice tests
- Assignments
- Access on mobile
- Full lifetime access

USE Coupon Code:
AWSYATRI066590703

→ 4 days left to get **FREE**
(Link is given in bio)

AWS AI Practitioner Exam Guide

- Yatharth Chauhan

Introduction to the AWS Certified AI (AIF-C01) Practitioner Exam

The AWS Certified AI Practitioner exam is a foundational-level certification designed to test your knowledge and understanding of artificial intelligence, machine learning, and generative AI technologies on the AWS platform. This certification is perfect for professionals who want to demonstrate their expertise in AI/ML without needing to be someone who codes or builds these systems from scratch.

Think of this certification as your way to show that you understand what AI and ML can do, how AWS services support AI/ML work, and how to use these technologies responsibly in your organization.

Who Should Take This Exam?

Target Candidate Profile

This exam is designed for individuals who have about 6 months of experience working with AI/ML technologies on AWS. You don't need to be someone who develops or builds these systems. Instead, you should be someone who uses these technologies and understands how they work.

Ideal Job Roles for This Certification

The exam is perfect for professionals in these types of roles:

- Business Analysts
- IT Support Specialists
- Marketing Professionals
- Product Managers or Project Managers

- Line-of-Business Managers or IT Managers
- Sales Professionals
- Anyone responsible for understanding or implementing AI/ML solutions in their organization

What Skills You Need Before Taking This Exam

Before you sit for this exam, you should already be comfortable with these AWS basics:

- **Core AWS Services:** You should know what Amazon EC2, Amazon S3, AWS Lambda, and Amazon SageMaker are used for
- **AWS Security Basics:** You should understand what the AWS shared responsibility model means
- **IAM (Identity and Access Management):** You should know how to secure and control access to AWS resources
- **AWS Infrastructure:** You should be familiar with concepts like AWS Regions, Availability Zones, and edge locations
- **AWS Pricing:** You should have a basic understanding of how AWS charges for services

Exam Format and Basic Information

Exam Overview

- **Exam Name:** AWS Certified AI Practitioner
- **Exam Code:** AIF-C01
- **Exam Duration:** 90 minutes
- **Total Questions:** 65 questions (50 scored + 15 unscored)
- **Passing Score:** 700 out of 1,000 (scaled score)
- **Cost:** USD 100
- **Format:** Available at Pearson VUE testing centers or online with a proctor

- **Languages Available:** 12 languages including English, Spanish, French, German, Japanese, Chinese, and others

Important Scoring Information

The exam uses a scaled score from 100 to 1,000, with a minimum passing score of 700. This scaling helps make sure that all exam versions are fair, even if some versions are slightly harder than others.

The exam also includes unscored questions that AWS uses to test questions before officially adding them to the exam. You won't know which questions are unscored, so treat all questions seriously.

Question Types You'll Encounter

The exam contains different types of questions to test your understanding in various ways:

Multiple Choice Questions

- You get one correct answer and three incorrect answers
- Select the one best answer
- This is the most common question type

Multiple Response Questions

- Two or more correct answers exist among five or more options
- You must select ALL correct answers to get credit
- Be careful here - partial credit is not given

Ordering Questions

- You get a list of 3 to 5 items that need to be arranged in the correct order
- You must place them in the correct sequence to receive credit
- These test your understanding of processes and workflows

Matching Questions

- You match items from one list to items in another list
- You get 3 to 7 prompts to match
- All pairs must be matched correctly to receive credit

Case Study Questions

- A scenario is presented with two or more questions about it

- The same scenario applies to each question
- Each question is scored separately, so you can get partial credit on case studies

Important Exam Rules

- Unanswered questions are marked as incorrect
- There is no penalty for guessing, so always provide an answer
- You don't need to pass each individual section - you just need to achieve an overall score of 700
- Section-level feedback is provided but should be used with caution when interpreting your performance

Certification Validity

Your AWS Certified AI Practitioner certification is valid for 3 years. Before it expires, you can:

- Retake the exam to recertify, or
 - Earn the AWS Certified Machine Learning Engineer - Associate certification, which automatically renews this certification
-

What the Exam Tests: The Five Domains

The exam covers five main areas of knowledge, each with a different weight or importance. Understanding these domains helps you know what to study and how much to focus on each area.

Domain 1: Fundamentals of AI and ML (20% of the exam)

This domain tests your basic understanding of artificial intelligence and machine learning concepts. About one-tenth of the questions on your exam will focus on this material.

Basic AI Concepts and Terminology You Need to Know

AI (Artificial Intelligence)

- The broad field of creating machines that can perform tasks that normally require human intelligence
- Includes machine learning, robotics, and other intelligent systems

- Example: A computer system that can understand and respond to human language

ML (Machine Learning)

- A subset of AI focused on systems that learn from data
- Instead of being explicitly programmed, ML systems improve through experience
- Example: A spam filter that learns what emails are spam by analyzing examples

Deep Learning

- A specialized type of machine learning using artificial neural networks
- Neural networks are inspired by how the human brain works
- Good for processing complex data like images and language
- Example: A system that can recognize faces in photographs

Neural Networks

- Computer systems inspired by biological neurons in the brain
- Made up of layers of interconnected nodes that process information
- Can learn complex patterns from data
- Example: A neural network trained to identify different types of animals in images

Computer Vision

- The field of AI that teaches computers to understand images and video
- Used for tasks like facial recognition, object detection, and medical imaging
- Example: A system that can identify defects in manufactured products by analyzing images

Natural Language Processing (NLP)

- The field of AI focused on understanding and working with human language
- Used for tasks like translation, sentiment analysis, and question answering
- Example: A chatbot that understands customer questions and provides helpful responses

Model

- The result of training a machine learning system
- A trained model can make predictions or decisions based on new data
- Example: After training on thousands of emails, a spam detection model can classify new emails as spam or not spam

Algorithm

- The step-by-step procedure that a machine learning system uses to learn from data
- Different algorithms are suited to different types of problems
- Example: Decision trees, random forests, and neural networks are different algorithms

Training

- The process of teaching a machine learning model using historical data
- The model learns patterns from this training data
- Example: Showing thousands of cat pictures to train a model to recognize cats

Inferencing

- The process of using a trained model to make predictions or decisions on new data
- Also called "prediction" or "inference"
- Happens after training is complete
- Example: Using a trained model to detect whether a new email is spam

Bias

- Systematic errors or unfairness in a machine learning system
- Can occur when training data doesn't represent all groups fairly
- Example: A hiring model trained mostly on successful male employees might be biased against women

Fairness

- The quality of a machine learning system treating all groups equally
- Related to reducing bias and ensuring ethical use
- Example: A loan approval system that fairly evaluates applications from all demographic groups

Fit

- How well a trained model performs on new data it hasn't seen before
- **Underfitting:** Model is too simple and doesn't learn patterns well (poor fit)
- **Overfitting:** Model memorizes training data too well but doesn't work on new data (also poor fit)
- **Good Fit:** Model learns patterns well and works on new data

Large Language Model (LLM)

- A type of neural network trained on massive amounts of text data

- Can understand and generate human-like text
- Powers systems like ChatGPT, Claude, and others
- Example: A model trained on billions of words that can answer questions and write essays

Understanding the Differences: AI vs ML vs Deep Learning

Think of these as nested categories:

- **AI** is the biggest umbrella - anything that makes machines intelligent
- **ML** is a major subset of AI - systems that learn from data
- **Deep Learning** is a subset of ML - uses neural networks to learn from complex data

All deep learning is machine learning, and all machine learning is AI, but not all AI is machine learning.

Types of Inferencing

Once your model is trained, you need to use it to make predictions. There are different ways to do this:

Batch Inferencing

- Process many predictions all at once in a batch
- Not real-time - there's a delay between when you request predictions and when you get them
- Used when you don't need immediate answers
- Example: Processing a company's entire customer list overnight to identify customers likely to leave

Real-Time Inferencing

- Get predictions immediately when you request them
- Lower latency - very fast response time
- Used when you need immediate answers
- Example: A website checking in real-time whether a transaction is fraudulent

Types of Data in Machine Learning

Machine learning models work with different types of data:

Labeled Data

- Data where the correct answer is already known
- Used for training supervised learning models
- Example: Past emails labeled as "spam" or "not spam"

Unlabeled Data

- Data without known answers
- Used for unsupervised learning to find hidden patterns
- Example: Customer transaction data without labels for "fraud" or "legitimate"

Tabular Data

- Data organized in rows and columns, like a spreadsheet
- Each row is a record and each column is a feature or characteristic
- Example: A table with columns for age, income, credit score, and whether a loan was defaulted

Time-Series Data

- Data collected at regular time intervals
- Values change over time
- Example: Stock prices recorded every minute, or temperature readings every hour

Image Data

- Pictures or visual information
- Can be used for computer vision tasks
- Example: Photos used to train a system to recognize different dog breeds

Text Data

- Written language or documents
- Used for natural language processing tasks
- Example: Customer reviews used to train a system to detect sentiment (positive or negative)

Structured Data

- Data that's organized in a defined format
- Follows a specific schema or structure
- Example: Customer database with defined fields for name, address, phone number

Unstructured Data

- Data without a specific format or structure
- Includes images, text, audio, and video
- Example: Social media posts or medical imaging scans

Three Main Types of Machine Learning

Supervised Learning

- You provide examples with correct answers during training
- The model learns to predict the right answer for new examples
- Used when you know what you're trying to predict
- Includes two sub-types:
 - **Regression:** Predicting continuous numbers (like predicting house prices)
 - **Classification:** Predicting categories (like predicting if an email is spam or not)
- Example: Training a model on past customer data to predict which customers will make a purchase

Unsupervised Learning

- You provide data without correct answers
- The model finds hidden patterns or groups in the data
- Used when you want to discover patterns you don't know about
- Main type: **Clustering** - finding natural groups in data
- Example: Analyzing customer purchase behavior to find groups of similar customers without being told what groups to look for

Reinforcement Learning

- A model learns through trial and error with rewards and penalties
- Similar to how humans learn - you get rewards for good actions and penalties for bad ones
- Used for complex decision-making scenarios
- Example: Training a computer to play chess by rewarding it when it wins and penalizing it when it loses

Task 1.1 Summary: Basic AI Concepts

You should be able to explain each of the above terms in simple language and understand when they are used. The exam will test whether you can recognize these concepts in real-world scenarios.

Task Statement 1.2: Identify Practical Use Cases for AI

Understanding when and where to use AI is just as important as understanding how AI works. This section teaches you to identify good opportunities for AI in business.

When AI and ML Provide Real Value

Assisting Human Decision Making

- AI can process huge amounts of information quickly
- Help humans make better decisions by providing insights and recommendations
- Example: A system that analyzes thousands of customers to help a bank decide who to approve for a loan

Scalability and Automation

- AI can automate repetitive tasks that would require many human workers
- Allows businesses to scale their operations without hiring proportionally more people
- Example: An automated system that can process 10,000 customer support tickets daily, whereas a human agent can only process 50

Pattern Recognition and Prediction

- AI excels at finding patterns in large datasets
- Can predict future trends or outcomes
- Example: Predicting which products a customer is likely to buy based on their past behavior

When NOT to Use AI/ML

It's crucial to recognize situations where AI/ML is not the right solution:

When Specific Outcomes Are Needed

- Sometimes you need a specific, guaranteed outcome, not a prediction
- AI gives probabilities, not certainties
- Example: If you need to calculate exact tax amounts, you need a rule-based system, not a machine learning model

When Cost Exceeds Benefits

- Building and maintaining AI systems is expensive
- Development costs, infrastructure costs, and ongoing monitoring costs add up
- If the cost of building the system exceeds the value it provides, don't do it

- Example: Building a complex AI system to solve a problem that affects only 5% of customers might not be worth it

When You Have Very Little Data

- ML models need sufficient training data to work well
- With too little data, the model won't learn proper patterns
- Traditional rule-based systems might work better

When Explainability Is Critical

- Some AI models (like neural networks) make predictions but it's hard to explain why
- In regulated industries like banking or healthcare, you might need to explain every decision
- Example: A bank might need to explain to a rejected loan applicant exactly why they were rejected - something an AI "black box" can't easily do

Selecting the Right ML Technique for Different Problems

Different types of business problems require different machine learning approaches:

Regression

- Used when you're predicting a continuous number
- The output is a number that can be any value within a range
- Examples:
 - Predicting house prices (could be any number)
 - Predicting temperature for tomorrow
 - Forecasting sales revenue
 - Predicting customer lifetime value

Classification

- Used when you're predicting a category or class
- The output is one of several predetermined categories
- Examples:
 - Is an email spam or not spam? (2 categories)
 - Is a customer likely to stay or leave? (2 categories)
 - What product category does an image belong to? (multiple categories)
 - Is a transaction fraudulent or legitimate?

Clustering

- Used when you want to find natural groups or segments in your data
- You don't know the groups in advance - you discover them
- Examples:
 - Grouping customers by similar behavior for targeted marketing
 - Finding clusters of diseases in medical data
 - Identifying groups of similar products
 - Discovering distinct user personas from behavioral data

Real-World AI Applications

Here are examples of how AI is actually being used today:

Computer Vision

- Facial recognition - identifying people in photos or videos
- Object detection - finding specific items in images (detecting defects in manufactured products)
- Medical imaging analysis - detecting diseases in X-rays or MRI scans
- Autonomous vehicles - understanding the environment around a car

Natural Language Processing

- Machine translation - translating text between languages
- Sentiment analysis - determining if text expresses positive, negative, or neutral emotion
- Question answering systems - answering questions in natural language
- Text summarization - automatically creating summaries of long documents

Speech Recognition

- Converting spoken words to text
- Powering virtual assistants that understand voice commands
- Transcribing meeting recordings
- Voice-based customer service systems

Recommendation Systems

- Recommending products to customers (like Netflix recommending movies)
- Suggesting friends on social media
- Recommending content based on user behavior
- Personalized shopping experiences

Fraud Detection

- Identifying suspicious credit card transactions
- Detecting insurance claim fraud
- Identifying unusual account access patterns
- Flagging suspicious login attempts

Forecasting

- Predicting future sales and demand
- Predicting equipment failure before it happens (predictive maintenance)
- Forecasting customer churn (who might leave)
- Predicting inventory needs

AWS Managed AI/ML Services You Should Know

AWS provides several ready-to-use AI services that you don't have to build yourself:

Amazon SageMaker

- AWS's main platform for machine learning
- Used for building, training, and deploying ML models
- Provides pre-built algorithms and templates
- Includes tools for data preparation and model monitoring

Amazon Transcribe

- Converts speech to text
- Can transcribe audio files and live audio streams
- Supports multiple languages
- Example: Transcribing customer service calls

Amazon Translate

- Translates text between languages
- Supports dozens of languages
- Uses neural machine translation
- Example: Translating customer feedback from Spanish to English

Amazon Comprehend

- Analyzes text to understand it
- Can detect the language, sentiment, key phrases, and entities in text
- Identifies topics in large collections of documents

- Example: Analyzing customer reviews to understand what customers like and dislike

Amazon Lex

- Builds conversational interfaces (chatbots)
- Understands user intent from text or speech
- Powers virtual assistants
- Example: A customer service chatbot that understands customer requests

Amazon Polly

- Converts text to speech
- Creates natural-sounding voice from written text
- Supports multiple languages and different voice options
- Example: Creating audio versions of documents or adding voice to applications

Task Statement 1.3: Describe the ML Development Lifecycle

Building a successful machine learning solution involves many stages.

Understanding this lifecycle helps you see how all the pieces fit together.

Components of an ML Pipeline

An ML pipeline is the entire process from raw data to a working model making predictions. Here are the major stages:

Data Collection

- Gathering raw data that will be used to train the model
- Data comes from databases, APIs, sensors, logs, or other sources
- The quality and quantity of data affects model quality
- Example: Collecting thousands of past customer transactions to train a fraud detection model

Exploratory Data Analysis (EDA)

- Understanding what your data looks like and what patterns it contains
- Visualizing data to see trends and relationships
- Identifying data quality issues like missing values or outliers
- Asking questions like: "What does the distribution of prices look like?" or "Are there any obvious patterns?"
- Important for deciding what to do next in the pipeline

Data Pre-processing

- Cleaning and preparing data for use in machine learning
- Handles tasks like:
- Removing duplicate records
- Handling missing values (filling them in or removing rows)
- Converting data to the right format
- Removing or fixing incorrect values
- Example: Converting all dates to the same format, removing customers with incomplete address information

Feature Engineering

- Creating new features (input variables) from raw data that help the model learn better
- Combines, transforms, or extracts information from raw data
- More art than science - requires domain knowledge
- Example: Creating a "age at purchase" feature by subtracting birth date from purchase date

Model Training

- The actual learning process where the model learns patterns from training data
- The model's parameters are adjusted to minimize errors on training data
- Different algorithms learn in different ways
- Takes time and computational resources
- Example: Training a neural network on 1 million customer records

Hyperparameter Tuning

- Fine-tuning the settings of how a model learns
- Different from model parameters - these are settings you choose before training
- Affects how well the model learns
- Example: Choosing the learning rate (how fast the model adjusts) or the number of layers in a neural network

Evaluation

- Testing how well the trained model performs
- Using separate test data that the model has never seen before
- Measuring accuracy and other performance metrics
- Determines if the model is good enough for real-world use

- Example: Testing the fraud detection model on 100,000 new transactions to see how accurate it is

Deployment

- Moving the trained model into production so it can make real predictions
- Setting up the infrastructure to run the model and serve predictions to applications
- Example: Deploying a fraud detection model so it can score new credit card transactions in real-time

Monitoring

- Continuously watching how the model performs after deployment
- Checking if the model's accuracy stays high over time
- Detecting if the model's performance degrades
- Deciding when to retrain the model with newer data

Sources of ML Models: Build vs. Use Pre-Trained

You have choices about where your model comes from:

Open Source Pre-Trained Models

- Models that have already been trained by others and are available for free
- Can often be used directly or fine-tuned for your specific task
- Examples: TensorFlow models, PyTorch models, Hugging Face models
- Advantages: Fast to get started, often high quality, free
- Disadvantages: Might not perfectly fit your specific use case

Custom-Built Models

- Training your own model from scratch using your own data
- Requires more time and expertise
- Can be tailored specifically to your business needs
- Better if you have unique data or special requirements
- Example: Training a model on your company's specific customer data

AWS SageMaker JumpStart

- AWS service providing pre-trained models ready to use
- Can be used as-is or fine-tuned for your needs
- Covers many common use cases
- Saves development time

Methods to Use a Model in Production

Once your model is trained, how do you actually use it? There are different approaches:

Managed API Service

- AWS hosts the model and provides an API you call to get predictions
- You send data to the API and receive predictions back
- AWS handles scaling, security, and maintenance
- Example: Amazon SageMaker endpoints that you call with data

Self-Hosted API

- You deploy the model on your own servers or containers
- You maintain the infrastructure yourself
- More control but more responsibility
- Example: Running a model in a Docker container on your own servers

Batch Processing

- Process many predictions all at once, not in real-time
- Good for less time-sensitive applications
- Example: Processing all customer records once per night to identify churn risks

AWS Services for Each ML Pipeline Stage

AWS provides specific services for different parts of the ML lifecycle:

Data Collection and Storage

- Amazon S3: Store raw data files
- AWS Glue: Catalog and organize your data
- Amazon Kinesis: Collect streaming data in real-time

Data Preparation

- Amazon SageMaker Data Wrangler: Clean and prepare data visually
- AWS Glue: Transform and clean data
- AWS Glue DataBrew: Clean data without coding

Feature Engineering

- Amazon SageMaker Feature Store: Store and manage features centrally

Model Training and Tuning

- Amazon SageMaker: Train models using built-in algorithms
- SageMaker Autopilot: Automatically find the best model for you

Model Evaluation and Monitoring

- Amazon SageMaker Model Monitor: Monitor model performance over time
- Amazon SageMaker Clarify: Detect bias in models

Deployment

- Amazon SageMaker: Deploy models as endpoints
- AWS Lambda: Serve predictions with serverless functions

Machine Learning Operations (MLOps)

MLOps is about making machine learning systems reliable and maintainable in production. Key concepts:

Experimentation

- Trying different approaches and comparing results
- Tracking which models work better
- Finding the best solution through systematic testing

Repeatable Processes

- Ensuring the same steps produce the same results every time
- Creating standard procedures everyone follows
- Makes it easier for teams to collaborate

Scalable Systems

- Building systems that can handle growing amounts of data and traffic
- Automatically adjusting resources based on demand
- Example: A system that can score 100 transactions per second without slowing down

Managing Technical Debt

- Avoiding shortcuts that create problems later
- Keeping code and models maintained and documented
- Example: Not taking shortcuts that make the code hard to understand later

Production Readiness

- Making sure the model is actually ready to use in production
- Passing quality checks and performance tests
- Having proper error handling and logging

Model Monitoring

- Continuously checking model performance in production

- Detecting when performance drops
- Setting up alerts when something goes wrong

Model Re-training

- Periodically updating the model with new data
- Keeping the model accurate as the world changes
- Deciding when to stop using an old model and use a new one

Performance Metrics: Technical and Business

To know if your model is working well, you need to measure it. There are technical metrics and business metrics:

Technical Performance Metrics

Accuracy

- Percentage of predictions that were correct
- Example: If the model is right 95% of the time, accuracy is 95%
- Good for general understanding but doesn't tell the whole story

Area Under the ROC Curve (AUC)

- Measures how well the model separates positive and negative cases
- Values range from 0.5 (random guessing) to 1.0 (perfect)
- Better for imbalanced datasets where one outcome is much more common

F1 Score

- Combines precision (accuracy of positive predictions) and recall (finding all positives)
- Useful when you care about both false positives and false negatives
- Example: In fraud detection, you care about finding fraud AND not falsely flagging legitimate transactions

Business Metrics

Cost Per User

- How much it costs to serve each user or process each transaction
- Example: If it costs \$1 to run the model for each customer prediction, that's your cost per user

Development Costs

- Total cost to build and deploy the model
- Includes data scientist time, infrastructure, tools, etc.

Customer Feedback

- What users and customers think about the system
- Are they satisfied with the results?
- Do they trust the system?

Return on Investment (ROI)

- The financial benefit of the model compared to its cost
 - Is the model making or saving the company money?
 - Example: If the fraud detection model saves \$1 million per month and costs \$10,000 per month to run, the ROI is high
-

Domain 1 Summary

Domain 1 focuses on the fundamental building blocks of AI and ML. You should be able to:

- Define and explain basic AI/ML terminology
 - Recognize when AI/ML is appropriate for a business problem
 - Understand the complete lifecycle of an ML project
 - Know which AWS services support different parts of the ML lifecycle
 - Understand both technical and business metrics for evaluating success
-

Domain 2: Fundamentals of Generative AI (24% of the exam)

Generative AI is a special type of AI that creates new content - text, images, code, and more. This is becoming increasingly important in business. About a quarter of your exam will focus on this area.

Task Statement 2.1: Explain Basic Concepts of Generative AI

Generative AI has its own terminology and concepts that are different from traditional machine learning. Let's break them down in simple terms.

Foundation Concepts of Generative AI

Tokens

- The basic units that language models work with
- Text is broken down into tokens - usually small pieces of words or whole words

- Example: The sentence "Hello, how are you?" might be broken into tokens like ["Hello", ",", "how", "are", "you", "?"]
- Models are charged by tokens, not by words - more tokens = higher cost
- Think of tokens like "words" but not exactly - sometimes punctuation is a token

Chunking

- Breaking large documents or texts into smaller, manageable pieces
- Each chunk is small enough to process efficiently
- Used when working with documents too large for the model to handle at once
- Example: Breaking a 50-page manual into 500 chunks of 100 tokens each

Embeddings

- Mathematical representations of text or images as numbers
- Convert words, sentences, or images into vectors (lists of numbers)
- Captures the meaning of the text in numerical form
- Things with similar meanings have similar embeddings
- Example: The words "cat" and "dog" have similar embeddings because they're both animals

Vectors

- Lists of numbers representing information
- Each number in the vector is a dimension
- Used in machine learning to represent complex concepts mathematically
- Example: A word might be represented as [0.2, 0.8, -0.5, 0.1, ...] with hundreds of numbers

Prompt Engineering

- The art of writing text instructions (prompts) to get the best responses from AI models
- Carefully crafting what you ask affects the quality of answers
- Different techniques produce different results
- Example: "Summarize this in 5 sentences" produces better results than "Summarize this"

Transformer-Based LLMs

- Large Language Models built using transformer architecture
- Transformers are the core technology behind models like ChatGPT and Claude
- They process all words in a sentence simultaneously (not one by one)

- This makes them very powerful and efficient
- Example: GPT-4, Claude, and others use transformer architecture

Foundation Models

- Very large models trained on enormous amounts of text, images, or other data
- Trained for general purposes, not for specific tasks
- Can be adapted for many different uses without much retraining
- Example: A model trained on all of Wikipedia and the internet becomes a foundation model

Multi-Modal Models

- Models that can work with multiple types of input - text, images, audio, or video
- Can understand relationships between different types of data
- Example: A model that can look at an image and describe what's in it, or create images from text descriptions

Diffusion Models

- A type of generative model that gradually creates images or other content
- Works by starting with random noise and gradually refining it into coherent content
- Used for image generation and other creative tasks
- Underlying technology in systems like DALL-E and Stable Diffusion

Potential Use Cases for Generative AI

Generative AI is good at creating new content. Here are common business applications:

Content Generation

- Writing articles, blogs, or marketing copy
- Creating product descriptions
- Generating social media posts
- Example: A company using AI to create personalized marketing emails for thousands of customers

Image and Video Generation

- Creating product images without photography
- Generating artwork or design elements
- Creating training videos
- Example: Generating multiple product variations to test which sells best

Code Generation

- Writing software code based on descriptions
- Helping developers write code faster
- Explaining and documenting code
- Example: Describing a function you want, and the AI writes the code

Summarization

- Creating short summaries of long documents
- Condensing meeting notes
- Summarizing articles or research papers
- Example: Automatically creating executive summaries of long reports

Translation

- Converting text from one language to another
- Maintaining meaning and nuance
- Example: Translating customer service chats in real-time

Chatbots and Virtual Assistants

- Creating conversational AI that understands and responds to customers
- Handling customer service inquiries
- Providing information and assistance
- Example: A customer service bot that can answer questions about products or troubleshoot issues

Search Enhancement

- Improving search results to be more relevant and natural
- Understanding user intent better
- Providing direct answers instead of just links
- Example: A search that provides an answer directly instead of just returning links

Recommendation Engines

- Suggesting products or content personalized to each user
- Based on understanding both the user and the items to recommend
- Example: Netflix recommending shows you might enjoy

The Foundation Model Lifecycle

Foundation models go through their own lifecycle, different from regular ML models:

Data Selection

- Choosing what data to train on
- Quality and diversity of data is critical
- Large, diverse datasets are important for foundation models
- Example: Deciding to include websites, books, news articles, and academic papers in training data

Model Selection

- Choosing the architecture - transformer or other types
- Deciding on model size
- Selecting pre-existing models vs. training from scratch
- Example: Deciding whether to train a new model or fine-tune an existing one like GPT-4

Pre-training

- Training the foundation model on large amounts of general data
- Takes enormous computational resources and time
- Creates the base model that understands language or images
- Usually only done by large companies like OpenAI, Google, Meta
- Example: Training a model on hundreds of billions of text tokens from the internet

Fine-tuning

- Adapting a pre-trained foundation model to be better for a specific task or domain
- Uses much less data and computation than pre-training
- Makes the model specialized for your needs
- Example: Taking a general-purpose model and fine-tuning it on medical texts to make it better at medical questions

Evaluation

- Testing the model to see if it works well
- Human evaluation - people judging if responses are good
- Automated evaluation using metrics and benchmarks
- Example: Having humans rate the quality of responses to test questions

Deployment

- Making the model available for use
- Creating APIs or interfaces people can use

- Deciding on access controls and safety measures
- Example: Deploying a model as a web interface customers can use

Feedback

- Collecting information on how well the model performs in real use
 - Learning from user feedback
 - Using feedback to improve the model
 - Example: Collecting examples of bad responses to improve the model
-

Task Statement 2.2: Understand Capabilities and Limitations of Generative AI

Generative AI is powerful but not perfect. You need to understand both what it can do well and its limitations.

Advantages of Generative AI

Adaptability

- Can handle many different tasks without retraining
- One model can be used for writing, answering questions, explaining things, and more
- Flexible and versatile
- Example: The same model that writes emails can also write code and answer questions

Responsiveness

- Can provide immediate responses to queries
- Much faster than waiting for a human expert
- Always available 24/7
- Example: Customer service chats provide instant responses instead of customers waiting for an email reply

Simplicity

- Easy to use - just write what you want
- No need to learn complex interfaces or programming
- Natural interaction using everyday language
- Example: Anyone can use ChatGPT without training - just start typing questions

Speed to Value

- Fast to implement and deploy
- Can see results quickly without lengthy development
- Reduces time from idea to deployment
- Example: A company can have a customer service chatbot running in days instead of months

Disadvantages and Limitations

Hallucinations

- The model generates incorrect information that sounds plausible
- Confidently states things that aren't true
- Can't distinguish between real knowledge and made-up information
- Very serious problem - users might believe false information
- Example: A model confidently explaining how to do something dangerous in a completely wrong way

Interpretability Issues

- Hard to understand why the model gave a particular response
- Doesn't explain its reasoning clearly
- Makes it difficult to trust or verify responses
- Example: A medical AI recommends a treatment but can't explain why

Inaccuracy

- Responses might be partially wrong or completely wrong
- Particularly bad with recent information, specialized domains, or complex reasoning
- Example: The model provides outdated information about current events

Non-Determinism

- Same question might get different answers each time you ask
- Makes behavior unpredictable
- Makes testing and validation difficult
- Example: Asking "What is 2+2?" might get "4" one time and "four" another time

Bias

- Models trained on biased data produce biased outputs
- Can perpetuate or amplify societal biases

- Example: A model trained on mostly male technical authors might suggest female developers are unsuitable for certain roles

Safety Concerns

- Models can produce harmful, illegal, or inappropriate content
 - Can be manipulated to bypass safety measures
 - Requires careful monitoring and controls
-

Task Statement 2.3: Describe AWS Infrastructure and Technologies for Generative AI

AWS provides several services and infrastructure specifically designed for generative AI. Understanding what's available helps you know what tools to use.

AWS Services for Generative AI

Amazon Bedrock

- Fully managed service for accessing foundation models
- Offers models from Amazon, Anthropic, Stability AI, and others
- No need to build or maintain your own model infrastructure
- Simple API to use models
- Example: Calling Amazon Bedrock to access Claude or Llama models without hosting them yourself

Amazon SageMaker JumpStart

- Pre-trained foundation models ready to use
- Can be deployed with just a few clicks
- Provides tutorials and sample notebooks
- Lower barrier to entry for generative AI
- Example: Finding a pre-trained image generation model and deploying it in minutes

PartyRock (Amazon Bedrock Playground)

- Visual interface for building generative AI applications
- No coding required
- Great for experimentation and prototyping
- Allows building chatbots and other applications without programming
- Example: A business analyst creating a customer service chatbot without writing code

Amazon Q

- Generative AI assistant for business and development
- Understands your company's documents and knowledge base
- Can answer questions about your specific business information
- Integrated with AWS services and other business tools
- Example: An employee asking "What's our cloud security policy?" and getting answers specific to your company's documents

Advantages of Using AWS Generative AI Services

Accessibility

- Easy to use even without deep AI expertise
- Pre-trained models available immediately
- Good for teams without specialized AI knowledge
- Example: Business teams can use generative AI without hiring AI specialists

Lower Barrier to Entry

- Don't need to build models from scratch
- Don't need to understand complex ML infrastructure
- Can start using generative AI quickly
- Example: A startup can use generative AI in weeks instead of months

Efficiency

- AWS handles scaling and optimization
- Services automatically scale up or down based on demand
- No need to manage infrastructure
- Example: Your application automatically handles 10x more traffic without you changing anything

Cost-Effectiveness

- Pay only for what you use
- Don't pay to maintain unused infrastructure
- More affordable than building in-house solutions
- Example: Using pay-per-token pricing instead of building your own model infrastructure

Speed to Market

- Deploy applications quickly
- Get feedback from users faster

- Iterate and improve quickly
- Example: Building and deploying a new feature in days instead of months

AWS Infrastructure Benefits for Generative AI

Security

- AWS provides enterprise-grade security
- Data is encrypted
- Access controls prevent unauthorized use
- Complies with security standards and regulations
- Example: Your proprietary data is protected with encryption and access controls

Compliance

- Meets regulatory requirements in various industries
- Audit trails for who accessed what
- Features like data residency requirements
- Example: Healthcare companies can ensure patient data stays in compliant regions

Responsibility

- AWS takes responsibility for infrastructure security
- You focus on your application, AWS secures the foundation
- Clear shared responsibility model
- Example: AWS secures the data center; you secure your API keys

Safety

- Built-in guardrails to prevent harmful outputs
- Features to detect and prevent problematic use
- Monitoring and controls
- Example: Tools to prevent the model from generating illegal content

Cost Considerations for AWS Generative AI Services

Different approaches have different cost tradeoffs:

Token-Based Pricing

- Pay per token used
- Input tokens and output tokens might have different prices
- Fine for variable workloads

- Becomes expensive with high volume
- Example: Processing 1 million input tokens and 500,000 output tokens

Provisioned Throughput

- Pay a fixed amount for guaranteed capacity
- Better for predictable, high-volume workloads
- More cost-effective for steady usage
- Example: Guaranteeing 100 API calls per second for a fixed monthly price

Response Time vs. Cost

- Faster models cost more
- Cheaper models might be slower
- Trade-off based on your needs
- Example: Using a smaller, cheaper model that's slower vs. a larger, faster model

Availability and Redundancy

- Higher availability costs more
- Redundancy across regions costs more
- Balance cost with reliability needs
- Example: Critical applications might use multiple regions; non-critical ones use a single region

Regional Coverage

- Services available in different AWS regions
- Some regions cost more than others
- Using regions close to your users is faster but might cost more
- Example: Using US East for cost savings vs. Asia Pacific for lower latency

Custom Models

- Building your own model costs more
- Requires more infrastructure and expertise
- Gives you more control and customization
- Example: Fine-tuning a model for your specific industry

Domain 2 Summary

Domain 2 focuses specifically on generative AI. You should understand:

- Key concepts like tokens, embeddings, prompt engineering, and foundation models

- What generative AI is good at and its limitations
 - Different use cases for generative AI
 - AWS services for building generative AI applications
 - The lifecycle of foundation models
 - Cost and performance tradeoffs
-

Domain 3: Applications of Foundation Models (28% of the exam)

This is the largest domain, covering how to actually build applications using foundation models. It's all about practical application.

Task Statement 3.1: Design Considerations for Foundation Model Applications

When building applications with foundation models, you need to make many design decisions. Understanding these choices helps you build better applications.

Selecting Pre-trained Models

When choosing a foundation model to use, consider:

Cost

- Different models have different pricing
- Larger models cost more per token
- High-volume applications need cost-effective models
- Example: A popular model might cost \$10 per 1 million input tokens, while a cheaper alternative costs \$2

Modality

- What type of input/output does the model support?
- Text-only, text with images, audio, video?
- Not all models support all types
- Example: Image generation requires a model with image generation capability; pure text models won't work

Latency

- How fast does the model respond?

- Some models are faster but less capable
- Real-time applications need fast models
- Example: A customer service chatbot needs fast responses; a report generator can take more time

Multi-Lingual Support

- Can the model work in your required languages?
- Not all models support all languages equally
- Important for global applications
- Example: If you serve Spanish customers, the model must be good at Spanish

Model Size

- Bigger models are generally more capable but slower and more expensive
- Smaller models are faster and cheaper but less capable
- Right choice depends on your needs
- Example: A large model for complex analysis; a small model for simple classification

Model Complexity

- More complex models have higher resource requirements
- More complex doesn't always mean better for your specific task
- Simpler models are easier to understand and debug
- Example: A simple model might work fine for sentiment analysis; you need complexity for multi-step reasoning

Customization

- Can you fine-tune or adapt the model?
- Some models allow customization; others don't
- Customization takes time and resources but improves results
- Example: Fine-tuning a model on your company's documents for better domain knowledge

Input/Output Length

- How much text can the model process at once?
- Different models have different limits
- Long documents need models with large context windows
- Example: Summarizing a 50-page document requires a model that can handle 50 pages of input

Inference Parameters and Their Effects

When using a model, you can adjust parameters that affect how it responds:

Temperature

- Controls how creative or random the model is
- Low temperature (0.0-0.3): Consistent, predictable, factual responses
- High temperature (0.7-1.0): Creative, varied, sometimes random responses
- Goldilocks zone (0.4-0.7): Balanced responses
- Example: Low temperature for generating SQL queries; high temperature for creative writing

Top-K

- Limits the model to choosing from the K most likely next words
- Helps prevent very unlikely words from being chosen
- Lower values make responses more consistent
- Example: Top-K of 40 means the model can only choose from the 40 most likely next words

Top-P

- Limits choices based on cumulative probability
- Model chooses from words that make up the top P percent of probability
- More natural than Top-K
- Example: Top-P of 0.9 means choosing from words until they add up to 90% of probability

Output Length

- How many tokens to generate
- Limit the length of responses
- Longer outputs cost more (more tokens)
- Example: Generating summaries by setting max output length to 100 tokens

Retrieval Augmented Generation (RAG)

RAG is a technique to make foundation models more accurate and relevant by providing them with specific information. This is crucial for many applications.

What is RAG?

- Retrieves relevant documents or data
- Provides that information to the foundation model

- The model uses that information to generate accurate responses
- Combines retrieval systems with generative models

How RAG Works

1. User asks a question
2. System finds relevant documents/information
3. Provides that information to the foundation model
4. Foundation model generates response based on the information provided
5. Response is more accurate and relevant

Business Applications of RAG

- Customer service using company knowledge bases
- Employee support systems using internal documentation
- Research tools that cite sources
- Domain-specific Q&A systems
- Example: A customer service chatbot that retrieves relevant product manuals before answering questions

Advantages of RAG

- Makes models more accurate for specific domains
- Reduces hallucinations (less made-up information)
- Can use current information the model wasn't trained on
- Provides source citations
- Example: A model can answer questions about your specific company policies even if it wasn't trained on them

Amazon Bedrock for RAG

- Amazon Bedrock Knowledge Base supports RAG
- Automatically retrieves relevant information for the model
- Simplifies building RAG applications
- Example: Set up a knowledge base with your documents; Bedrock handles retrieval automatically

Vector Databases for Embeddings

Foundation models work with embeddings, which are stored in vector databases. Understanding this is important for RAG and other applications.

What is a Vector Database?

- Stores embeddings (numerical representations of text or images)

- Allows searching for similar embeddings quickly
- Enables semantic search - finding things by meaning, not just keywords
- Example: Storing embeddings of all your company documents for quick retrieval

AWS Services for Vector Storage

Amazon OpenSearch Service

- Open-source search and analytics engine
- Can store and search vectors
- Scalable and reliable
- Supports semantic search on text

Amazon Aurora

- Relational database with vector support
- Can store both traditional data and vectors together
- Good if you already use Aurora
- Combines structured data with semantic search

Amazon Neptune

- Graph database that can store and query vectors
- Good for highly connected data
- Useful for relationship-heavy applications

Amazon DocumentDB

- MongoDB-compatible database
- Supports vector search
- Good for document storage with semantic search

Amazon RDS for PostgreSQL

- Relational database with vector extensions
- Familiar for traditional database users
- Good mix of structured data and vector search

Cost Tradeoffs for Customization

Different ways to customize foundation models have different cost implications:

Pre-training

- Training a foundation model from scratch
- Extremely expensive
- Only for very large organizations with massive compute budgets

- Rarely necessary
- Example: Only companies like OpenAI, Google, Meta do this

Fine-tuning

- Adapting a pre-trained model to your domain
- Expensive but much less than pre-training
- Requires time and computational resources
- Improves accuracy for your specific use case
- Example: Fine-tuning a model on medical texts to make it better at medical questions
- Cost: Thousands of dollars typically

In-Context Learning

- Providing examples in the prompt to teach the model
- No additional training needed
- Cheaper than fine-tuning
- Works reasonably well with foundation models
- Example: Showing the model a few examples of the format you want, then asking it to do the same for new data
- Cost: Just the tokens used; no training cost

RAG (Retrieval Augmented Generation)

- Providing relevant information to the model
- No model training needed
- Cheapest approach for domain-specific accuracy
- Good for keeping information current
- Example: Retrieving relevant documents before asking the model to answer
- Cost: Just API calls for retrieval and generation

Ranking: Cheapest to Most Expensive

1. In-context learning (just API calls)
2. RAG (API calls plus retrieval)
3. Fine-tuning (training cost plus API calls)
4. Pre-training (massive cost, rarely done)

Agents for Multi-Step Tasks

Sometimes you need the model to perform complex tasks that require multiple steps. Agents help with this.

What are Agents?

- Autonomous systems that can break down complex tasks into steps
- Plan what to do
- Execute actions
- Check results
- Adjust as needed

Agents for Amazon Bedrock

- AWS service for building autonomous AI agents
- Models decide what actions to take
- Can integrate with other AWS services and tools
- Example: An agent that can check your calendar, read emails, and schedule meetings automatically

How Agents Work

1. Receive a complex task
2. Break it into steps
3. Decide what action to take
4. Take that action (call a tool or service)
5. Observe the result
6. Decide next action based on result
7. Repeat until task is complete

Real-World Examples

- Customer service agents that can check order status, process returns, and update shipping
- Research agents that search information, analyze data, and write reports
- Travel booking agents that check flights, hotels, and prices
- Code writing agents that write, test, and debug code

Task Statement 3.2: Choose Effective Prompt Engineering Techniques

Prompt engineering is the art and science of writing instructions to get the best results from foundation models. This is increasingly important because better prompts = better results.

Concepts and Constructs of Prompt Engineering

Context

- Background information that helps the model understand your request
- Tells the model who it is, what it's doing, and background information
- Example: "You are a medical doctor. A patient asks: [question]"

Instruction

- Clear command of what you want the model to do
- Should be specific and unambiguous
- Example: "Summarize this in 5 sentences" vs. "Tell me about this"

Negative Prompts

- Explicitly telling the model what NOT to do
- Helps prevent unwanted outputs
- Example: "Do not include personal opinions" or "Do not use technical jargon"

Examples

- Showing the model examples of what you want
- Often more effective than just describing it
- Example: Showing three examples of good customer service responses, then asking the model to write a similar response

Model Latent Space

- The internal representation the model uses
- Different prompts navigate this space differently
- Some prompts land in better parts of the space
- Technical concept but important for understanding why different prompts produce different results

Techniques for Better Prompts

Chain-of-Thought

- Ask the model to show its reasoning step by step
- Results in better, more accurate responses
- Particularly good for complex problems
- Example: Instead of "What's the answer?", ask "Show me how you would solve this step by step"
- Result: Better answers with explanations

Zero-Shot

- Asking the model to do something with no examples

- Works okay for simple tasks
- Less reliable for complex tasks
- Example: "Classify this email as spam or not spam" without giving examples

One-Shot

- Providing one example of what you want
- Works better than zero-shot
- Sometimes not enough for complex tasks
- Example: Showing one example of a bad tweet, then asking it to identify other bad tweets

Few-Shot

- Providing multiple examples (usually 3-5)
- Usually works best for complex tasks
- More tokens (costs more) but better results
- Example: Showing 5 examples of customer complaints and how they should be categorized

Prompt Templates

- Reusable prompt structures for common tasks
- Save time and ensure consistency
- Can be filled in with different information each time
- Example: A template for "Summarize this document" that you can use over and over

Benefits and Best Practices

Response Quality Improvement

- Good prompts get significantly better results
- Small changes in phrasing can make big differences
- Well-written prompts save you having to fix bad outputs

Experimentation

- Try different approaches
- Test what works best
- Continuous improvement
- Example: Testing whether chain-of-thought or zero-shot works better for your use case

Guardrails

- Setting boundaries on responses
- Preventing harmful outputs
- Protecting your brand
- Example: "Do not mention competitors" or "Keep responses professional"

Discovery

- Learning what the model is capable of
- Finding unexpected uses
- Testing limits
- Example: Discovering the model can generate code as well as text

Specificity and Concision

- Be specific about what you want
- Keep prompts concise
- Avoid unnecessary words
- Example: "Write a professional email requesting a meeting" is better than "Write an email about a meeting"

Using Comments or Sections

- Break complex prompts into sections
- Use comments to explain what each section does
- Makes prompts easier to understand and modify
- Example: Separating the context section, instruction section, and examples section

Risks and Limitations of Prompt Engineering

Exposure

- Sensitive information might be revealed in prompts
- Proprietary information could be leaked
- Data in prompts might be used to train future models
- Example: Don't include confidential company secrets in prompts

Poisoning

- Deliberately crafted inputs designed to break the model
- Can cause the model to behave unexpectedly
- Security risk

- Example: An attacker crafting a prompt to make the model generate harmful content

Hijacking

- Taking control of the model's behavior through prompts
- Overriding intended behavior
- Security concern
- Example: An attacker's prompt overriding a chatbot's safety guidelines

Jailbreaking

- Bypassing safety measures
 - Getting the model to do things it's supposed to refuse
 - Ethical concern
 - Example: Tricking the model into generating illegal content
-

Task Statement 3.3: Describe Training and Fine-Tuning of Foundation Models

Sometimes you need to adapt foundation models for your specific needs. Understanding how this works is important.

Key Elements of Training Foundation Models

Pre-training

- Initial training on massive amounts of general data
- Done by model creators (OpenAI, Anthropic, etc.)
- Teaches the model general language or visual understanding
- Takes enormous computational resources
- Takes weeks or months
- Rarely something you'll do

Fine-tuning

- Adapting a pre-trained model to be better for your specific task or domain
- Uses much smaller amounts of data than pre-training
- Takes hours or days
- Much more affordable
- Something you'll do if you want domain specialization
- Example: Taking a general model and teaching it about your industry

Continuous Pre-training

- Keeping the model up-to-date with new information
- Additional pre-training on newer data
- Used when you need the model to know about recent events or changes
- Example: Retraining a model on new customer interactions to improve performance

Methods for Fine-tuning Foundation Models

Instruction Tuning

- Training the model to follow instructions better
- Showing it examples of instructions and good responses
- Improves instruction-following ability
- Example: Training a model on customer service instructions so it behaves better as a support agent

Adapting for Specific Domains

- Specializing the model for your industry or domain
- Using domain-specific data in training
- Results in better performance on domain tasks
- Example: Training a medical model on medical texts so it's better at medical questions

Transfer Learning

- Using knowledge learned from one task to help with another task
- Saves training time and data
- Common approach in ML
- Example: A model trained on general text that's fine-tuned on medical text

Continuous Pre-training

- Updating the model with new data as it becomes available
- Keeps model knowledge current
- Example: Updating a model with the latest product information

Data Preparation for Fine-tuning

Data Curation

- Carefully selecting training data
- Quality over quantity

- Relevant to your specific needs
- Example: Selecting only high-quality customer interactions for training

Governance

- Ensuring data is properly managed
- Tracking data sources
- Ensuring data quality and compliance
- Example: Making sure all data is properly licensed and doesn't contain personal information

Data Size

- Enough data to teach the model
- More data generally leads to better results
- But diminishing returns - 1000 examples might be 90% as good as 10000
- Depends on complexity of task
- Example: 500 examples might be enough for simple tasks; complex tasks might need thousands

Data Labeling

- Marking up data to show what's correct
- Input-output pairs showing what the model should learn
- Labor-intensive but important
- Example: Labeling examples as "good customer service response" or "bad customer service response"

Representativeness

- Data should represent the real situations the model will face
- Diverse data prevents bias
- Example: Training data should include interactions from different customer types and regions

Reinforcement Learning from Human Feedback (RLHF)

- Using human feedback to improve the model
- Humans rate responses as good or bad
- Model learns to generate responses humans prefer
- Used by many modern models
- Example: Humans rating model-generated responses, then retraining based on those ratings

Task Statement 3.4: Evaluate Foundation Model Performance

After training or fine-tuning a model, you need to evaluate how well it works.

Approaches to Evaluation

Human Evaluation

- People read and judge model outputs
- The most reliable evaluation method
- Time-consuming and expensive
- Often the best way to know if users will be satisfied
- Example: Having experts read and rate 100 model responses to judge quality

Benchmark Datasets

- Standardized test datasets with known correct answers
- Allows comparison across models
- Objective and repeatable
- Example: Using a standard dataset to test how well a model answers questions

Metrics for Assessing Foundation Model Performance

ROUGE (Recall-Oriented Understudy for Gisting Evaluation)

- Measures quality of summarization
- Compares system summary to human-created summaries
- Scores range from 0-1, higher is better
- ROUGE-1 looks at individual words
- ROUGE-2 looks at two-word phrases
- ROUGE-L looks at longest matching sequences
- Example: Testing a summarization system to see how well its summaries match human summaries

BLEU (Bilingual Evaluation Understudy)

- Measures quality of machine translation
- Compares machine translation to human translations
- Scores range from 0-1, higher is better
- Looks at overlapping words and phrases
- Example: Evaluating a Spanish-to-English translation system

BERTScore

- Modern metric using deep learning

- Compares semantic similarity using embeddings
- Better at catching meaning preservation
- Good for many NLG (Natural Language Generation) tasks
- Example: Evaluating whether a generated paraphrase preserves the meaning of the original

Determining if Model Meets Business Objectives

It's not enough for a model to be technically good - it needs to meet business goals:

Productivity

- Does the model help people work faster?
- Reduces time on tasks?
- Increases output?
- Example: Does a code-writing model help developers write code faster?

User Engagement

- Do users actually use the system?
- Are they satisfied with results?
- Do they come back?
- Example: Do customers engage with a chatbot, or do they abandon it?

Task Efficiency

- Does the system accomplish tasks efficiently?
- Cost per task?
- Time per task?
- Example: Does the automated system process customer requests cheaper than humans?

ROI (Return on Investment)

- Is the system making money for the company?
- Cost vs. benefit?
- Example: Does the automation save more money than it costs?

Domain 3 Summary

Domain 3 is about applying foundation models in real situations. You should understand:

- How to select and design applications using foundation models
 - Prompt engineering techniques to get better results
 - How to fine-tune models for specific needs
 - How to evaluate whether models actually work
 - Cost tradeoffs of different approaches
 - Practical tools like RAG and vector databases
-

Domain 4: Guidelines for Responsible AI (14% of the exam)

Building AI systems responsibly is increasingly important. This domain covers ethical considerations and risks.

Task Statement 4.1: Develop Responsible AI Systems

Responsible AI means building systems that are fair, safe, and trustworthy. This isn't just ethics - it's good business.

Features of Responsible AI

Bias

- Systematic unfairness in predictions or decisions
- Can harm individuals or groups
- Example: A hiring system that unfairly rejects women candidates
- Need to detect and reduce bias

Fairness

- Treating all groups equally
- No discrimination
- Everyone gets fair treatment regardless of protected characteristics
- Example: Loan approval system that evaluates applications fairly regardless of race or gender

Inclusivity

- Including diverse perspectives in design and training
- Ensuring the system works for all groups
- Not just optimizing for the majority

- Example: Making sure a chatbot works equally well for all languages and cultures

Robustness

- System works reliably in different situations
- Handles unusual inputs without breaking
- Resists adversarial attacks
- Example: A model that still works when given inputs slightly different from training data

Safety

- System doesn't cause harm
- Prevents dangerous, illegal, or unethical outputs
- Has guardrails and controls
- Example: A chatbot that refuses to help with illegal activities

Veracity

- Information provided is accurate and truthful
- Doesn't spread misinformation
- Citations when making claims
- Example: A system that tells you when it's uncertain instead of making up answers

Using Tools to Identify Responsible AI Features

Guardrails for Amazon Bedrock

- AWS tool to implement safety controls
- Prevents models from generating harmful content
- Can filter inputs and outputs
- Sets boundaries on what the model can say
- Example: Blocking requests that ask the model to generate illegal content

Responsible Practices for Model Selection

Environmental Considerations

- Training models uses enormous energy
- Larger models use more energy
- Consider environmental impact

- Example: Using smaller models when they work well enough, not always choosing the largest

Sustainability

- Building for long-term use
- Not constantly retraining unnecessarily
- Efficient infrastructure
- Example: Designing systems that work well for years, not months

Legal Risks of Working with Generative AI

Intellectual Property Infringement

- Models trained on copyrighted material without permission
- Generated content might violate copyrights
- Legal liability
- Example: Model trained on copyrighted books might reproduce passages

Biased Model Outputs

- Model generates discriminatory content
- Company could be sued
- Reputational damage
- Example: Hiring system that unfairly rejects candidates from certain groups

Loss of Customer Trust

- If AI system causes harm, customers lose trust
- Affects business
- Can take years to rebuild
- Example: Medical AI giving wrong advice erodes trust in the company

End User Risk

- Users rely on AI and might be harmed
- Users might trust AI too much
- Company has responsibility for consequences
- Example: Autonomous vehicle giving unsafe navigation instructions

Hallucinations

- AI making up false information
- Users believe false information
- Can cause serious harm

- Legal and reputational consequences
- Example: Medical AI confidently giving wrong treatment advice

Dataset Characteristics for Responsible AI

Good datasets are crucial for responsible AI:

Inclusivity

- Represents diverse populations
- Not just majority groups
- Example: Training data for facial recognition should include diverse ethnicities

Diversity

- Wide range of examples
- Covers different scenarios
- Prevents overfitting to specific cases
- Example: Training a chatbot on conversations from many different regions and demographics

Curated Data Sources

- Data comes from quality sources
- Verified and vetted
- Not random internet data
- Example: Using professionally collected data instead of unverified web scraping

Balanced Datasets

- Not dominated by one class or group
- Equal representation when possible
- Example: Equal numbers of approved and denied loan applications in training data

Effects of Bias and Variance

Bias (Systematic Error)

- Consistent error in one direction
- Model systematically wrong
- Can affect demographic groups differently
- Example: A model always predicting slightly too low, or consistently underestimating women's qualifications

Variance (Inconsistency)

- Different results from similar inputs
- Inconsistent predictions
- Overfitting to training data
- Example: Same qualifications getting approved sometimes and denied other times

Effects on Demographic Groups

- Bias and variance affect different groups differently
- Some groups might be disproportionately harmed
- Example: A model with bias against a particular gender affects that gender most

Inaccuracy

- Wrong predictions
- Bad decisions based on wrong information
- Example: Predictions are wrong 30% of the time

Overfitting

- Model memorizes training data
- Works great on training data but poorly on new data
- Lacks generalization
- Example: Model trained on one company's data doesn't work for other companies

Underfitting

- Model too simple
- Doesn't capture patterns in data
- Works poorly even on training data
- Example: Linear model trying to fit a complex non-linear relationship

Tools to Detect and Monitor Bias

Analyzing Label Quality

- Are the labels (correct answers) in training data actually correct?
- Bad labels lead to bad models
- Need to verify label quality
- Example: Checking if the "spam" and "not spam" labels in email training data are accurate

Human Audits

- Having people review model decisions
- Catch problems humans can spot
- Provides feedback for improvement
- Example: Auditors reviewing a sample of model decisions to check for bias

Subgroup Analysis

- Analyzing how model performs on different subgroups
- Do some groups get worse accuracy?
- Identifies discrimination
- Example: Checking if a hiring model treats men and women equally

Amazon SageMaker Clarify

- AWS tool for detecting bias
- Analyzes models for bias
- Identifies which groups are affected
- Provides explanations for decisions
- Example: Automatically detecting bias in a loan approval model

SageMaker Model Monitor

- Monitors model performance over time
- Detects performance degradation
- Tracks data quality
- Alerts when problems occur
- Example: Detecting that model accuracy has dropped from 95% to 92%

Amazon Augmented AI (Amazon A2I)

- Human-in-the-loop tool
- Combines human review with automation
- Humans review uncertain or high-stakes predictions
- Improves accuracy and catches problems
- Example: For high-value loans, humans review model recommendations before approval

Task Statement 4.2: Transparent and Explainable Models

Users need to understand why an AI system made a particular decision. This builds trust and is sometimes legally required.

Transparent vs. Opaque Models

Transparent Models

- You can understand how they work
- Decisions are explainable
- Internal logic is clear
- Example: Decision trees - you can see exactly what rules led to a decision

Explainable Models

- Might be complex internally but can be explained
- Provide explanations for decisions
- Why did you get this result?
- Example: A model can show which factors most influenced its decision

Opaque or Black-Box Models

- You can't easily understand why they make predictions
- Neural networks are often opaque
- Hard to explain decisions
- Example: A deep learning model makes good predictions but it's unclear why

Tools for Transparent and Explainable Models

Amazon SageMaker Model Cards

- Document models and their characteristics
- Transparency about model capabilities and limitations
- Document training data and evaluation results
- Helps stakeholders understand the model
- Example: Model card documenting a hiring model's accuracy on different demographic groups

Open Source Models

- Code is publicly available
- Can be inspected and understood
- Community can audit for problems
- More transparent than proprietary models
- Example: Open-source models from Hugging Face that anyone can inspect

Data and Licensing

- Clear documentation of what data was used
- Understanding potential biases

- License compliance
- Source attribution
- Example: Documenting that a model was trained on Wikipedia, public news sources, and academic papers

Tradeoffs Between Safety and Transparency

Measuring Interpretability

- How explainable is the model?
- Quantifying how much humans can understand decisions
- Some models are more interpretable than others
- Example: Decision tree (very interpretable) vs. deep neural network (less interpretable)

Performance

- More interpretable models might be less accurate
- Less interpretable models might be more accurate
- Balance based on needs
- Example: A simple decision tree is less accurate but very interpretable; a complex neural network is more accurate but less interpretable

Human-Centered Design for Explainable AI

Explaining to Users

- Explanations should be understandable to users
- Not technical ML explanations
- Show factors that matter to the user
- Example: Instead of technical feature importance, show "Your application was rejected because income was below threshold"

Trust Building

- Transparent AI builds trust
- Users need to understand and believe in the system
- Example: Showing users exactly why a recommendation was made builds confidence

Accountability

- Someone is responsible for decisions
- Audit trail of how decisions were made

- Ability to appeal or contest
 - Example: A system that records and explains every decision for later review
-

Domain 4 Summary

Domain 4 focuses on responsible and trustworthy AI. You should understand:

- What responsible AI means and why it matters
 - Common risks and legal issues
 - Tools for detecting and preventing bias
 - The importance of transparency and explainability
 - How to evaluate AI systems for fairness
-

Domain 5: Security, Compliance, and Governance (14% of the exam)

The last domain covers protecting AI systems from security threats and ensuring compliance with regulations.

Task Statement 5.1: Explain Methods to Secure AI Systems

Securing AI systems is important to protect data and prevent misuse.

AWS Services and Features for Securing AI

IAM Roles, Policies, and Permissions

- Control who can access what
- Principle of least privilege - only give access needed
- Audit trail of who did what
- Example: Giving a data scientist access to training data but not production systems

Encryption

- Protect data so only authorized people can read it
- Encryption in transit (while moving) and at rest (while stored)
- Uses cryptographic keys
- Example: Encrypting sensitive customer data with AES-256

Amazon Macie

- Finds sensitive data automatically

- Scans storage for personal information
- Alerts when sensitive data is found
- Helps prevent data breaches
- Example: Finding credit card numbers or health records stored unprotected

AWS PrivateLink

- Private connection to AWS services
- Data doesn't go over the internet
- Very secure
- Example: Accessing SageMaker training through a private connection

AWS Shared Responsibility Model

- AWS secures the infrastructure
- You secure your application and data
- Clear division of responsibility
- Example: AWS secures the data centers; you secure your passwords and API keys

Source Citation and Data Lineage

Data Lineage

- Tracking where data came from
- How it was transformed
- Where it went
- Provides audit trail
- Example: Tracking that customer data came from the database, was cleaned, then used for training

Data Cataloging

- Maintaining inventory of data
- Metadata about what data exists
- Who owns what data
- Example: A catalog showing all datasets in the company and who has access

SageMaker Model Cards

- Document model origins
- Training data sources
- Data transformations

- Traceability
- Example: Model card showing what data was used to train the model and when

Best Practices for Secure Data Engineering

Assessing Data Quality

- Is the data clean and accurate?
- Are there errors or inconsistencies?
- Quality problems lead to bad results
- Example: Checking for duplicate records or missing values

Privacy-Enhancing Technologies

- Techniques to protect privacy
- Differential privacy - adding noise to protect individuals
- Federated learning - training without centralizing data
- Data anonymization - removing identifying information
- Example: Removing names and contact info from customer data before analysis

Data Access Control

- Only authorized people see data
- Different levels of access for different roles
- Audit log of access
- Example: Marketing team can see aggregate data but not individual customer records

Data Integrity

- Data hasn't been modified
- Detection of corruption or tampering
- Checksums and verification
- Example: Ensuring training data hasn't been accidentally modified

Security and Privacy Considerations

Application Security

- Securing the AI application itself
- Protection against attacks
- Code review and testing
- Example: Protecting the API that serves model predictions

Threat Detection

- Identifying attack attempts
- Monitoring for suspicious activity
- Alerts when problems detected
- Example: Detecting unusual API access patterns

Vulnerability Management

- Finding and fixing security weaknesses
- Regular testing
- Patching known vulnerabilities
- Example: Running security scans to find weak spots

Infrastructure Protection

- Securing the underlying systems
- Firewalls and network security
- Physical security of data centers
- Example: AWS securing data centers with multiple layers of protection

Prompt Injection

- Attackers inserting malicious text into prompts
- Can hijack model behavior
- Security risk
- Example: An attacker's text in a prompt causing the model to ignore its normal guidelines

Encryption at Rest and in Transit

- Protecting data when stored (at rest)
- Protecting data when being sent (in transit)
- Uses encryption keys
- Example: Data encrypted on disk and encrypted when sent over the network

Task Statement 5.2: Recognize Governance and Compliance

Organizations need to follow regulations and have policies for AI systems.

Regulatory Compliance Standards

ISO (International Organization for Standardization)

- International standards for quality and security

- ISO 27001 for information security
- ISO standards for AI systems emerging
- Demonstrates commitment to standards
- Example: Being ISO 27001 certified shows your security practices meet international standards

SOC (System and Organization Controls)

- Controls for security, availability, and confidentiality
- SOC 2 for cloud services
- Third-party audit of controls
- Example: Being SOC 2 Type II certified shows you maintain security controls

Algorithm Accountability Laws

- Laws requiring transparency in algorithms
- Some jurisdictions require explaining algorithmic decisions
- Right to appeal algorithmic decisions
- Example: GDPR in Europe giving people the right to explanation of algorithmic decisions

AWS Services for Governance and Compliance

AWS Config

- Tracks configuration of resources
- Compliance rules checking
- Configuration history
- Identifies non-compliance
- Example: Tracking that all S3 buckets have encryption enabled

Amazon Inspector

- Scans for security vulnerabilities
- Automated security assessment
- Identifies weaknesses
- Example: Scanning for unpatched software or open security groups

AWS Audit Manager

- Helps with compliance audits
- Evidence collection automated
- Audit trail
- Example: Collecting evidence that security controls are in place

AWS Artifact

- Provides compliance documents
- Can download compliance reports
- Shows AWS certifications
- Example: Downloading SOC 2 reports for your own audit

AWS CloudTrail

- Logs all API calls
- Who did what and when
- Audit trail
- Compliance evidence
- Example: CloudTrail recording that an API key was created at 3 PM on Tuesday

AWS Trusted Advisor

- Best practices recommendations
- Cost optimization suggestions
- Security recommendations
- Example: Recommending closing unused databases or enabling MFA

Data Governance Strategies

Data Lifecycles

- What data do you have?
- How long do you keep it?
- When do you delete it?
- Compliance with regulations
- Example: Keeping customer data for 7 years then deleting it per regulations

Logging

- Recording who accessed what data
- When was it accessed?
- What was done with it?
- For audit purposes
- Example: Logging every access to customer personal data

Data Residency

- Where is data physically stored?
- Some regulations require data to stay in certain countries

- Data sovereignty issues
- Example: GDPR requires European data to stay in Europe

Monitoring and Observation

- Continuous checking of systems
- Detecting anomalies
- Early warning of problems
- Example: Monitoring that training accuracy is stable over time

Data Retention

- How long to keep data
- Based on legal requirements and business needs
- Balance between keeping data and storage costs
- Example: Deciding to keep transaction logs for 7 years

Governance Processes and Frameworks

Policies

- Rules about how AI systems are used
- Who can build AI systems?
- What gets reviewed before deployment?
- Documentation
- Example: Policy that all AI systems must be audited for bias before deployment

Review Cadence

- How often are AI systems reviewed?
- Regular ongoing reviews
- Triggered by changes or problems
- Example: Reviewing model performance monthly

Review Strategies

- How are systems reviewed?
- Manual human review
- Automated testing
- External audits
- Example: Both automated tests and human experts reviewing the system

Governance Frameworks

- Structured approach to governance

- Clear process and responsibilities
- Example: Using a framework like the Generative AI Security Scoping Matrix

Generative AI Security Scoping Matrix

- AWS framework for managing generative AI risks
- Helps identify risks and controls
- Example: Systematic approach to evaluating generative AI systems for security

Transparency Standards

- Being transparent about how AI works
- Documenting model capabilities and limitations
- Public disclosure of AI usage
- Example: Company policy to explain to customers when AI is making decisions about them

Team Training Requirements

- People building AI systems need training
- Understanding responsible AI
- Understanding security requirements
- Keeping knowledge current
- Example: All team members doing annual training on responsible AI practices

Domain 5 Summary

Domain 5 covers protecting and governing AI systems. You should understand:

- Security tools and best practices
- Compliance standards and regulations
- Governance frameworks and processes
- Data protection and privacy
- The role of compliance tools and services

Summary of All Domains

The AWS Certified AI Practitioner exam tests your knowledge across five domains:

Domain 1: Fundamentals of AI and ML (20%)

Basic concepts, identifying use cases, understanding the ML lifecycle

Domain 2: Fundamentals of Generative AI (24%)

Generative AI concepts, capabilities and limitations, AWS services for generative AI

Domain 3: Applications of Foundation Models (28%)

Designing applications, prompt engineering, fine-tuning, evaluation

Domain 4: Guidelines for Responsible AI (14%)

Responsible AI practices, detecting bias, transparency and explainability

Domain 5: Security, Compliance, and Governance (14%)

Security practices, compliance standards, governance frameworks

Exam Preparation Tips

Study Strategy

1. **Start with the basics:** Understand Domain 1 thoroughly before moving to other domains
2. **Build progressively:** Move from fundamentals to applications
3. **Use real examples:** Think about how concepts apply to real businesses
4. **Practice with case studies:** The exam includes case studies - practice analyzing scenarios
5. **Focus on relationships:** Understand how domains relate to each other
6. **Learn AWS services:** Know what each AWS service does and when to use it

Question Type Strategies

Multiple Choice

- Read all options before answering
- Eliminate obviously wrong answers
- Look for the best answer, not just a correct answer

Multiple Response

- You must get ALL correct answers for credit
- Be careful - it's easy to select most but miss one
- If unsure, re-read the question to verify you found all correct answers

Ordering

- Think about the logical sequence
- Understand why one step comes before another
- Test your order by walking through it

Matching

- Look for obvious pairs first
- Process of elimination helps with harder pairs
- Verify each match makes sense

Case Study

- Read the scenario thoroughly
- Each question is independent
- Use information from the scenario to answer

Time Management

- You have 90 minutes for 65 questions
- That's about 1.4 minutes per question
- Don't get stuck - move on and come back
- Leave time to review at the end

What to Study

- Know the AWS services mentioned in the exam guide
- Understand the concepts, not just definitions
- Learn why something matters, not just what it is
- Be able to apply concepts to scenarios
- Understand tradeoffs and when to choose different approaches

Before Exam Day

1. Review the official exam guide (which you have)
2. Take practice tests to identify weak areas
3. Study AWS documentation for services you're weak on

4. Review your notes the day before
5. Get good sleep the night before
6. Arrive early to the testing center

On Exam Day

1. Read questions carefully - many have subtle differences
 2. Don't rush - you have enough time
 3. Flag difficult questions and come back to them
 4. On multiple response questions, verify you got all correct answers
 5. Trust your knowledge - don't overthink
 6. Remember that some questions are unscored, so don't worry if something seems off
-

Quick Reference: AWS Services for AI/ML

In-Scope Services (What You Need to Know)

Core ML Services

- Amazon SageMaker: Main ML platform
- Amazon Bedrock: Managed foundation models
- Amazon Q: Business intelligence AI

Specific AI Services

- Amazon Comprehend: Text analysis
- Amazon Transcribe: Speech to text
- Amazon Translate: Language translation
- Amazon Polly: Text to speech
- Amazon Lex: Conversational interfaces
- Amazon Rekognition: Image recognition
- Amazon Textract: Extract text from images
- Amazon Personalize: Recommendations
- Amazon Kendra: Enterprise search
- Amazon Fraud Detector: Fraud detection

Data Services

- Amazon S3: Store data
- AWS Glue: Data preparation and cataloging
- AWS Glue DataBrew: Data cleaning without coding
- Amazon EMR: Big data processing
- AWS Lake Formation: Data lake management
- Amazon OpenSearch: Search and analytics
- Amazon Redshift: Data warehouse

Database Services (for storing embeddings and models)

- Amazon Aurora: Relational database
- Amazon DynamoDB: NoSQL database
- Amazon Neptune: Graph database
- Amazon DocumentDB: MongoDB-compatible
- Amazon RDS for PostgreSQL: Relational database
- Amazon ElastiCache: In-memory cache

Security and Governance

- AWS IAM: Access control
- AWS KMS: Encryption key management
- Amazon Macie: Data discovery and protection
- AWS CloudTrail: Audit logging
- AWS Artifact: Compliance documents
- AWS Audit Manager: Compliance audits

Monitoring and Analytics

- Amazon CloudWatch: Monitoring
- AWS Config: Configuration management
- Amazon Inspector: Security assessment
- AWS Trusted Advisor: Best practices

Final Thoughts

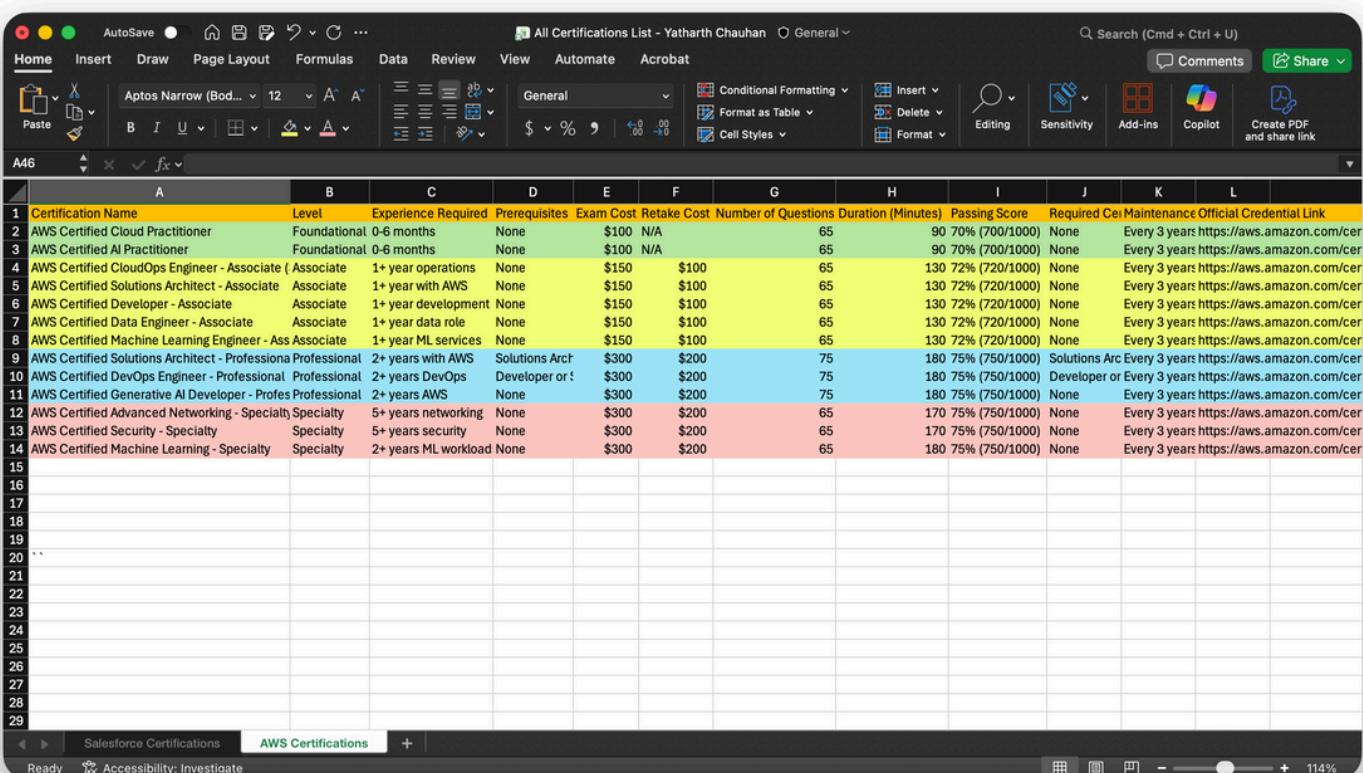
The AWS Certified AI Practitioner exam tests whether you understand AI and ML concepts and how they apply on AWS. You don't need to know how to code or build AI systems - you need to know what's possible, when to use AI, and how to use it responsibly.

Study the concepts, understand the use cases, and familiarize yourself with AWS services. Take practice tests, and you'll be well-prepared for exam day.

Remember: This certification is about showing that you can have intelligent conversations about AI/ML with your team and make good decisions about using these technologies. Focus on understanding, not memorization.

Good luck on your exam!

All AWS Certifications (Total: 13)



1	Certification Name	Level	Experience Required	Prerequisites	Exam Cost	Retake Cost	Number of Questions	Duration (Minutes)	Passing Score	Required Cer	Maintenance	Official Credential	Link
2	AWS Certified Cloud Practitioner	Foundational	0-6 months	None	\$100	N/A	65	90	70% (700/1000)	None	Every 3 years	https://aws.amazon.com/cer	
3	AWS Certified AI Practitioner	Foundational	0-6 months	None	\$100	N/A	65	90	70% (700/1000)	None	Every 3 years	https://aws.amazon.com/cer	
4	AWS Certified CloudOps Engineer - Associate	(Associate	1+ year operations	None	\$150	\$100	65	130	72% (720/1000)	None	Every 3 years	https://aws.amazon.com/cer	
5	AWS Certified Solutions Architect - Associate	Associate	1+ year with AWS	None	\$150	\$100	65	130	72% (720/1000)	None	Every 3 years	https://aws.amazon.com/cer	
6	AWS Certified Developer - Associate	Associate	1+ year development	None	\$150	\$100	65	130	72% (720/1000)	None	Every 3 years	https://aws.amazon.com/cer	
7	AWS Certified Data Engineer - Associate	Associate	1+ year data role	None	\$150	\$100	65	130	72% (720/1000)	None	Every 3 years	https://aws.amazon.com/cer	
8	AWS Certified Machine Learning Engineer - Ass Associate	Associate	1+ year ML services	None	\$150	\$100	65	130	72% (720/1000)	None	Every 3 years	https://aws.amazon.com/cer	
9	AWS Certified Solutions Architect - Professional	Professional	2+ years with AWS	Solutions Arch	\$300	\$200	75	180	75% (750/1000)	Solutions Arc	Every 3 years	https://aws.amazon.com/cer	
10	AWS Certified DevOps Engineer - Professional	Professional	2+ years DevOps	Developer or S	\$300	\$200	75	180	75% (750/1000)	Developer or	Every 3 years	https://aws.amazon.com/cer	
11	AWS Certified Generative AI Developer - Profes Professional	Professional	2+ years AWS	None	\$300	\$200	75	180	75% (750/1000)	None	Every 3 years	https://aws.amazon.com/cer	
12	AWS Certified Advanced Networking - Specialty	Specialty	5+ years networking	None	\$300	\$200	65	170	75% (750/1000)	None	Every 3 years	https://aws.amazon.com/cer	
13	AWS Certified Security - Specialty	Specialty	5+ years security	None	\$300	\$200	65	170	75% (750/1000)	None	Every 3 years	https://aws.amazon.com/cer	
14	AWS Certified Machine Learning - Specialty	Specialty	2+ years ML workload	None	\$300	\$200	65	180	75% (750/1000)	None	Every 3 years	https://aws.amazon.com/cer	
15													
16													
17													
18													
19													
20													
21													
22													
23													
24													
25													
26													
27													
28													
29													

Comment **#awsyatri** to get sheet in your DM

→ Connect is Required to
Enable DM



Follow me on
LinkedIn

→ to stay updated

Follow Now